

A PARSER-GENERATOR IN 100₁₆ LINES OF C++

Please excuse the longish intro – I promise this is going somewhere!

In the past, my day job consisted of creating and maintaining a Material Flow Control System (often called MFCS_{opt}) for warehouses and production plants. This necessitated connecting to various PLCs controlling the mechanical parts – from huge cranes, through conveyor belts, and all the way down to LED systems telling humans what to do.

All those systems had one particular thing in common: they all defined similar, but different text-based protocols to be used to communicate with them over TCP. In a simplified example, that’s how a message to a crane could be defined:

Field name	Field type	Comment
Begin	ALPHA[1]	Character [
Message Type	ALPHA[3]	“MOV” for move
X To	NUM[3]	
Y To	NUM[3]	
End	ALPHA[1]	Character]

move crane

Have you ever needed to implement a third-party plaintext protocol? It’s as simple as it’s boring. And Deity forbid if the documentation changes after initial implementation. You’ll waste so much time! At least that’s what I told my boss when I started creating a templated declarative parser.

To be fair, I was fairly accurate. I inherited code that used `std::map<std::string, std::string>`, and I wager that I wasted multiple days hunting all the typos in those strings.

Since C++ is a fairly strongly-typed language, there is no need for that – we should be able to leverage the type system to ensure that both our keys and values are correct. Let’s discuss the API:

- keys (field names) should be verified at compilation time – none of these pesky typos can pass here,
- values need to be of correct type, not the all-catching `std::string`,
- the code should be as close as possible to the documentation. Ideally, it’d be the documentation.

For example, we could want our MOV telegram to be defined as follows:

```
using mov = message<
    element<struct begin, char_constant<[' '>>,
    element<struct message_type, text<3>>,
    element<struct x_to, number<3>>,
    element<struct y_to, number<3>>,
    element<struct end, char_constant<'] '>>
>;
```

The usage should be also simple. For receiving:

```
auto data = socket.read();
mov m = mov::parse(data);
log << m.value<x_to>() << m.value<y_to>();
```

And for sending:

```
mov m;
m.value<message_type>() = "MOV";
m.value<x_to>() = 13;
m.value<y_to>() = 37;
socket.send(m.to_string());
```

This approach is Good Enough™. We have type safety, and we can even extend it to use custom types. For example, the above will write the following to the socket (note the padding: zeros for numbers, text would use spaces):

```
[MOV013037]
```

Moving on, the internal implementation is surprisingly simple. The main class template accepts a list of key-type pairs as variadic pack. It uses keys only to map them to values. The type has a bit more to do – each type is expected to know its length, and how to serialize and deserialize itself (or signal an error).

```
template<typename... Elements>
struct message
{
    static message parse(string_view buf);
    void write(char* buf, size_t size) const;
    string to_string() const;
    template<typename Key>
    constexpr auto& value();
private:
    tuple<typename d::element_value_type<
        Elements>::type...> data;
};
```

Class message definition – shortened and modified to fit here

```
template<size_t Length>
struct number
{
    static constexpr size_t length = Length;
    using value_type =
        d::type_to_hold_number<Length>;
    static void write(value_type const& val,
        char* buf);
    static value_type parse(string_view buf);
};
```

Class number definition – shortened and modified to fit here

As of writing this article, the whole `proto.hpp` has 248 lines, and I haven’t performed any line-saving optimizations on the file.

The code may be accessed at the following address: https://github.com/KrzaQ/protocol_parser_generator.